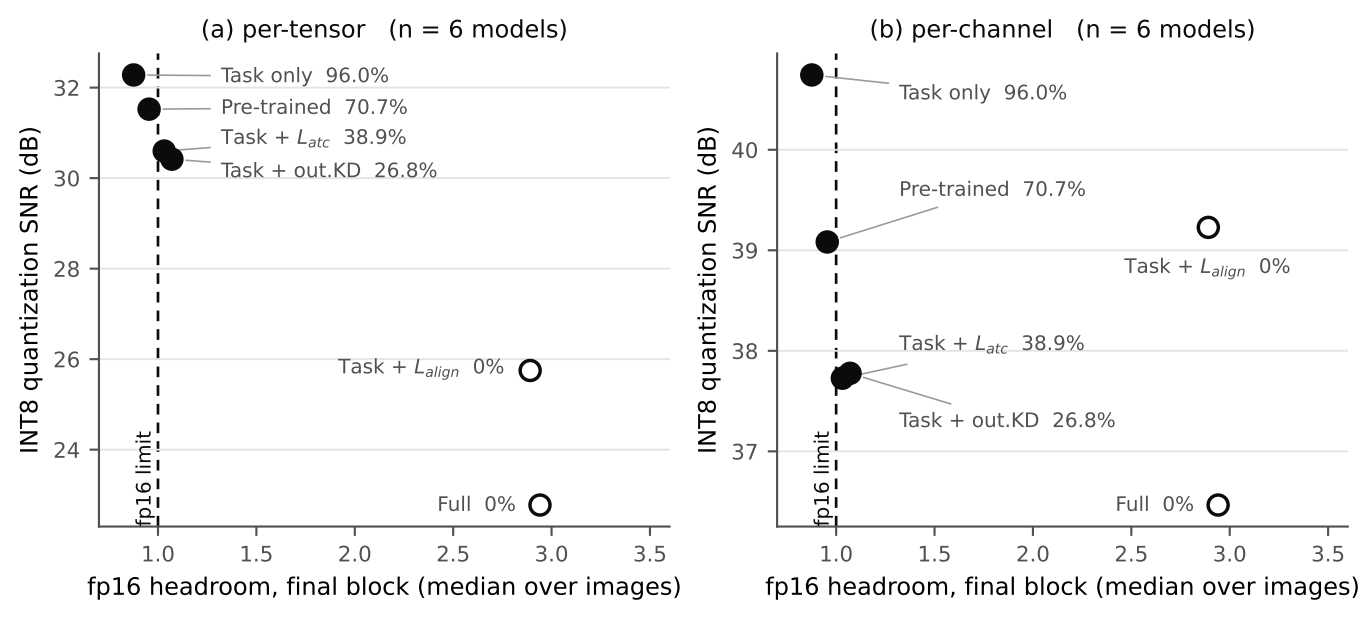
(a) per-tensor (n = 6 models)
INT8 quantization SNR (dB)
fp16 headroom, final block (median over images)
fp16 limit
Task only 96.0%
Pre-trained 70.7%
Task + L_{atc} 38.9%
Task + out.KD 26.8%
Task + L_{align} 0%
Full 0%
(b) per-channel (n = 6 models)
INT8 quantization SNR (dB)
fp16 headroom, final block (median over images)
fp16 limit
Task only 96.0%
Pre-trained 70.7%
Task + L_{align} 0%
Task + L_{atc} 38.9%
Task + out.KD 26.8%
Full 0%