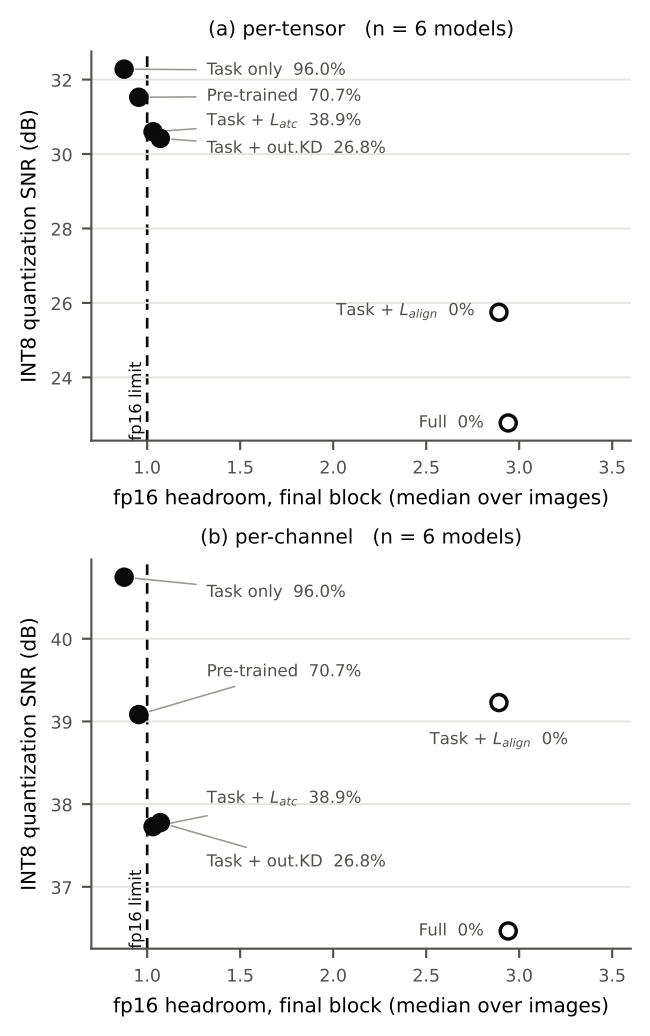
(a) per-tensor (n = 6 models)
INT8 quantization SNR (dB)
32
30
28
26
24
Task only 96.0%
Pre-trained 70.7%
Task + L_{atc} 38.9%
Task + out.KD 26.8%
Task + L_{align} 0%
Full 0%
fp16 limit
1.0
1.5
2.0
2.5
3.0
3.5
fp16 headroom, final block (median over images)
(b) per-channel (n = 6 models)
INT8 quantization SNR (dB)
40
39
38
37
Task only 96.0%
Pre-trained 70.7%
Task + L_{align} 0%
Task + L_{atc} 38.9%
Task + out.KD 26.8%
Full 0%
fp16 limit
1.0
1.5
2.0
2.5
3.0
3.5
fp16 headroom, final block (median over images)