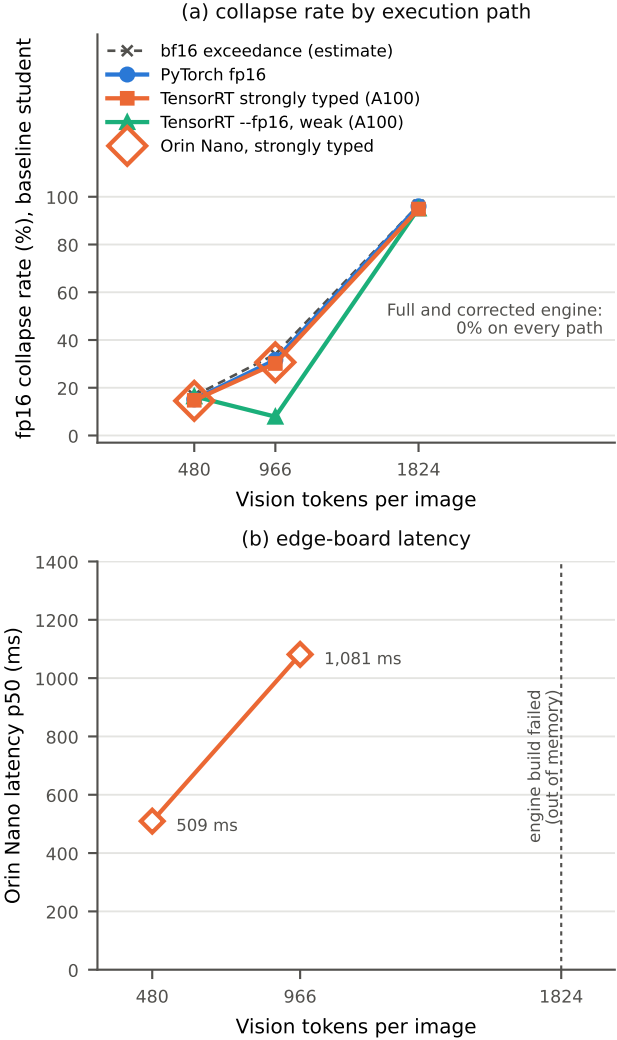
(a) collapse rate by execution path
bf16 exceedance (estimate)
PyTorch fp16
TensorRT strongly typed (A100)
TensorRT --fp16, weak (A100)
Orin Nano, strongly typed
fp16 collapse rate (%), baseline student
100
80
60
40
20
0
Full and corrected engine:
0% on every path
480
966
1824
Vision tokens per image
(b) edge-board latency
Orin Nano latency p50 (ms)
1400
1200
1000
800
600
400
200
0
1,081 ms
509 ms
engine build failed
(out of memory)
480
966
1824
Vision tokens per image